



Mawlana Bhashani Science & Technology University

Lab Report

Report No: 03

Course code: ICT-3208

Course title: Network Planning and Designing Lab

Date of Performance:

Date of Submission:

Submitted by

Name: Md. Rafatul Haque
ID:IT-17037
3rd year 2nd semester
Session: 2016-2017
Dept. of ICT
MBSTU.

Submitted To

Nazrul Islam
Assistant Professor
Dept. of ICT
MBSTU.

Lab Report Name: SDN Controllers and Mininet.

Objective:

- Install and use traffic generators as powerful tools for testing network performance.
- Install and configure SDN Controller.
- Install and understand how the mininet simulator works.
- Implement and run basic examples for understanding the role of the controller and how it interact with mininet.

Theory:

Traffic Generator:

What is iPerf?: iPerf is a tool for active measurements of the maximum achievable bandwidth on IP networks. It supports tuning of various parameters related to timing, buffers and protocols (TCP, UDP, SCTP with IPv4 and IPv6). For each test it reports the bandwidth, loss, and other parameters.

Controller:

OVS-testcontroller is a simple OpenFlow controller that manages any number of switches over the OpenFlow protocol, causing them to function as L2 MAC-learning switches or hubs. It is suitable for initial testing of OpenFlow networks.

Ryu is a component-based software defined networking framework. Ryu provides software components with well-defined API that make it easy for developers to create new network management and control applications. Ryu supports various protocols for managing network devices, such as OpenFlow, Netconf, OF-config, etc. About OpenFlow, Ryu supports fully 1.0, 1.2, 1.3, 1.4, 1.5 and Nicira Extensions. All of the code is freely available under the Apache 2.0 license.

Mininet: Mininet creates a realistic virtual network, running real kernel, switch and application code, on a single machine (VM, cloud or native) Because you can easily interact with your network using the Mininet CLI (and API), customize it, share it with others, or deploy it on real hardware, Mininet is useful for development, teaching, and research. Mininet is also a great way to develop, share, and experiment with OpenFlow and Software-Defined Networking systems.

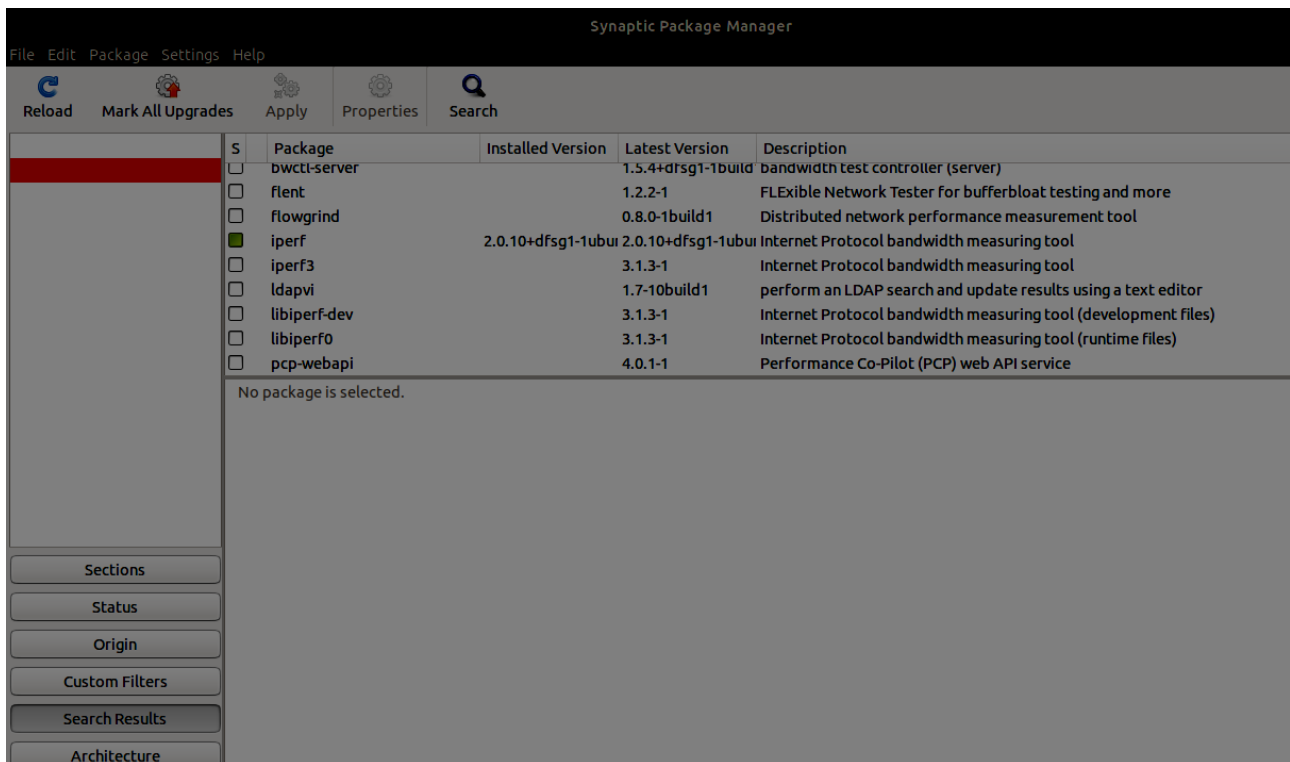
Methodology

TIP: For getting extra space in your USB-please use the following tips:

- Empty the trash
- Delete the Android related stuff
- Delete the extras for other courses
- Delete the already installed package sources

Install iperf

1. Open the Synaptic Package Manager (Navigator ->System-> Synaptic Package Manager)
2. Setup the proxy:
 - ❖ Click on settings-> Preference -> Network
 - ❖ Click on manual proxy configuration
 - ❖ HTTP and FTP Proxy: proxy.rmit.edu.au Port: 8080
3. Search for Quick filter `iperf`
4. Click on Mark for installation
5. Then click on Apply and wait until the package is installed



Install mininet

1. In Synaptic Package
2. Search for Quick filter `mininet`
3. Click on Mark for installation
4. Then click on Apply and wait until the package is installed

Install Controller

OVS controller:

1. In Synaptic Package
2. Search for Quick filter `openvswitch-controller`
3. Click on Mark for installation
4. Then click on Apply and wait until the package is installed

Part 1: Everyday Mininet Usage

Display Startup Options

Let's get started with Mininet's startup options.

Type the following command to display a help message describing Mininet's startup options:

\$ sudo mn -h

```
rafatul@rafatul-HP-Notebook:~$ sudo mn -h
[sudo] password for rafatul:
Usage: mn [options]
(type mn -h for details)

The mn utility creates Mininet network from the command line. It can create
parametrized topologies, invoke the Mininet CLI, and run tests.

Options:
  -h, --help                show this help message and exit
  --switch=SWITCH            default|ivs|lxb|ovs|ovsbr|ovsk|user[,param=value...]
                             ovs=OVSSwitch default=OVSSwitch ovsk=OVSSwitch
                             lxb=LinuxBridge user=UserSwitch ivs=IVSSwitch
                             ovsbr=OVSBridge
  --host=HOST                cfs|proc|rt[,param=value...]
                             rt=CPULimitedHost{'sched': 'rt'} proc=Host
                             cfs=CPULimitedHost{'sched': 'cfs'}
  --controller=CONTROLLER  default|none|nox|ovsc|ref|remote|ryu[,param=value...]
                             ovsc=OVSController none=NullController
                             remote=RemoteController default=DefaultController
                             nox=NOX ryu=Ryu ref=Controller
  --link=LINK                default|ovs|tc|tcu[,param=value...] default=Link
                             ovs=OVSLink tcu=TCULink tc=TCLink
  --topo=TOPO                linear|minimal|reversed|single|torus|tree[,param=value
                             ...] linear=LinearTopo torus=TorusTopo tree=TreeTopo
                             single=SingleSwitchTopo
                             reversed=SingleSwitchReversedTopo minimal=MinimalTopo
  -c, --clean                clean and exit
  --custom=CUSTOM            read custom classes or params from .py file(s)
  --test=TESTS               cli|build|pingall|pingpair|iperf|all|iperfudp|none|pin
                             gpair|iperfudp|pingall|iperfUDP
  -x, --xterms               spawn xterms for each node
  -i IPBASE, --ipbase=IPBASE
                             base IP address for hosts
  --mac                       automatically set host MACs
  --arp                       set all-pairs ARP entries
  -v VERBOSITY, --verbosity=VERBOSITY
```

```
                             info|warning|critical|error|debug|output
  --innamespace              sw and ctrl in namespace?
  --listenport=LISTENPORT   base port for passive switch listening
                             don't use passive listening port
  --nolistenport            don't use passive listening port
  --pre=PRE                  CLI script to run before tests
  --post=POST                CLI script to run after tests
  --pin                      pin hosts to CPU cores (requires --host cfs or --host
                             rt)
  --nat                      [option=val...] adds a NAT to the topology that
                             connects Mininet hosts to the physical network.
                             Warning: This may route any traffic on the machine
                             that uses Mininet's IP subnet into the Mininet
                             network. If you need to change Mininet's IP subnet,
                             see the --ipbase option.
  --version                  prints the version and exits
  --cluster=server1,server2...
                             run on multiple servers (experimental!)
  --placement=block|random  node placement for --cluster (experimental!)
```

Start Wireshark

To view control traffic using the OpenFlow Wireshark dissector, first open wireshark in the background:

\$ sudo wireshark &

```
rafatul@rafatul-HP-Notebook:~$ sudo wireshark &  
[1] 10191  
rafatul@rafatul-HP-Notebook:~$ QStandardPaths: XDG_RUNTIME_DIR not set, defaulting  
to '/tmp/runtime-root'
```

In the Wireshark filter box, enter this filter, then click **Apply**:

of

In Wireshark, click Capture, then Interfaces, then select Start on the loopback interface (lo).

For now, there should be no OpenFlow packets displayed in the main window.

Interact with Hosts and Switches

Start a minimal topology and enter the CLI:

\$ sudo mn

```
rafatul@rafatul-HP-Notebook:~$ sudo mn  
*** Creating network  
*** Adding controller  
*** Adding hosts:  
h1 h2  
*** Adding switches:  
s1  
*** Adding links:  
(h1, s1) (h2, s1)  
*** Configuring hosts  
h1 h2  
*** Starting controller  
c0  
*** Starting 1 switches  
s1 ...  
*** Starting CLI:  
mininet>
```

If no specific test is passed as a parameter, the Mininet CLI comes up.

In the Wireshark window, you should see the kernel switch connect to the reference controller.

Display Mininet CLI commands:

mininet>help

```
mininet> help

Documented commands (type help <topic>):
=====
EOF      gterm  iperfudp  nodes      pingpair    py      switch
dpctl    help   link      noecho     pingpairfull  quit    time
dump     intfs  links     pingall    ports       sh      x
exit     iperf  net       pingallfull px          source  xterm

You may also send a command to a node using:
  <node> command {args}
For example:
  mininet> h1 ifconfig

The interpreter automatically substitutes IP addresses
for node names when a node is the first arg, so commands
like
  mininet> h2 ping h3
should work.

Some character-oriented interactive commands require
noecho:
  mininet> noecho h2 vi foo.py
However, starting up an xterm/gterm is generally better:
  mininet> xterm h2
```

Display nodes:

mininet>nodes

```
mininet> nodes
available nodes are:
c0 h1 h2 s1
```

Display links:

mininet>net

```
mininet> net
h1 h1-eth0:s1-eth1
h2 h2-eth0:s1-eth2
s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0
c0
```

Dump information about all nodes:

mininet>dump

```
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=12056>
<Host h2: h2-eth0:10.0.0.2 pid=12058>
<OVSSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None pid=12063>
<OVSController c0: 127.0.0.1:6653 pid=12049>
```

You should see the switch and two hosts listed.

If the first string typed into the Mininet CLI is a host, switch or controller name, the command is executed on that node. Run a command on a host process:

mininet> h1 ifconfig -a

```
mininet> h1 ifconfig -a
h1-eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.1 netmask 255.0.0.0 broadcast 10.255.255.255
    inet6 fe80::84df:20ff:fe69:2d9a prefixlen 64 scopeid 0x20<link>
    ether 86:df:20:69:2d:9a txqueuelen 1000 (Ethernet)
    RX packets 47 bytes 6049 (6.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 16 bytes 1216 (1.2 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

You should see the host's **h1-eth0** and loopback (**lo**) interfaces. Note that this interface (**h1-eth0**) is not seen by the primary Linux system when **ifconfig** is run, because it is specific to the network namespace of the host process.

In contrast, the switch by default runs in the root network namespace, so running a command on the “switch” is the same as running it from a regular terminal:


```
mininet> s1 ifconfig -a
enp1s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    ether 30:e1:71:92:19:b9 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 9062 bytes 700457 (700.4 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 9062 bytes 700457 (700.4 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ovs-system: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether 1a:ef:5a:36:dd:ca txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether 7a:77:14:a8:9c:4a txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 20 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::5085:e0ff:fe29:44a8 prefixlen 64 scopeid 0x20<link>
    ether 52:85:e0:29:44:a8 txqueuelen 1000 (Ethernet)
    RX packets 16 bytes 1216 (1.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
```

mininet>s1 ifconfig -a

```

s1: flags=4098<BROADCAST,MULTICAST> mtu 1500
    ether 7a:77:14:a8:9c:4a txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 20 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::5085:e0ff:fe29:44a8 prefixlen 64 scopeid 0x20<link>
    ether 52:85:e0:29:44:a8 txqueuelen 1000 (Ethernet)
    RX packets 16 bytes 1216 (1.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 47 bytes 6049 (6.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

s1-eth2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::fc0f:67ff:feb1:12d6 prefixlen 64 scopeid 0x20<link>
    ether fe:0f:67:b1:12:d6 txqueuelen 1000 (Ethernet)
    RX packets 16 bytes 1216 (1.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 47 bytes 6069 (6.0 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

wlp2s0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::a5ea:3d39:2263:3213 prefixlen 64 scopeid 0x20<link>
    ether 30:e3:7a:4e:64:70 txqueuelen 1000 (Ethernet)
    RX packets 209382 bytes 262490176 (262.4 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 157891 bytes 21959411 (21.9 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

This will show the switch interfaces, plus the VM's connection out (**eth0**).

For other examples highlighting that the hosts have isolated network state, run **arp** and **route** on both **s1** and **h1**.

It would be possible to place every host, switch and controller in its own isolated network namespace, but there's no real advantage to doing so, unless you want to replicate a complex multiple-controller network. Mininet does support this; see the **-innamespace** option.

Note that *only* the network is virtualized; each host process sees the same set of processes and directories. For example, print the process list from a host process:

```
mininet> h1 ps -a
```

```
mininet> h1 ps -a
  PID TTY          TIME CMD
 1190 tty1      00:00:00 gnome-session-b
 1256 tty1      00:00:23 gnome-shell
 1289 tty1      00:00:00 Xwayland
 1331 tty1      00:00:00 ibus-daemon
 1334 tty1      00:00:00 ibus-dconf
 1336 tty1      00:00:00 ibus-x11
 1396 tty1      00:00:00 gsd-xsettings
 1399 tty1      00:00:00 gsd-a11y-settin
 1400 tty1      00:00:00 gsd-clipboard
 1401 tty1      00:00:01 gsd-color
 1404 tty1      00:00:00 gsd-datetime
 1405 tty1      00:00:00 gsd-housekeepin
 1406 tty1      00:00:00 gsd-keyboard
 1410 tty1      00:00:00 gsd-media-keys
 1411 tty1      00:00:00 gsd-mouse
 1412 tty1      00:00:00 gsd-power
 1413 tty1      00:00:00 gsd-print-notif
 1416 tty1      00:00:00 gsd-rfkill
 1421 tty1      00:00:00 gsd-screensaver
 1424 tty1      00:00:00 gsd-sharing
 1428 tty1      00:00:00 gsd-smartcard
 1431 tty1      00:00:00 gsd-sound
 1433 tty1      00:00:00 gsd-wacom
 1447 tty1      00:00:00 ibus-engine-sim
 1628 tty2      00:06:43 Xorg
 1645 tty2      00:00:00 gnome-session-b
 1784 tty2      00:08:49 gnome-shell
 1830 tty2      00:00:04 ibus-daemon
 1834 tty2      00:00:00 ibus-dconf
 1838 tty2      00:00:00 ibus-x11
 1911 tty2      00:00:00 gsd-power
 1912 tty2      00:00:00 gsd-print-notif
 1915 tty2      00:00:00 gsd-rfkill
 1919 tty2      00:00:00 gsd-screensaver
 1922 tty2      00:00:01 gsd-sharing
```

```
1919 tty2 00:00:00 gsd-screensaver
1922 tty2 00:00:01 gsd-sharing
1928 tty2 00:00:00 gsd-smartcard
1930 tty2 00:00:00 gsd-sound
1936 tty2 00:00:00 gsd-xsettings
1937 tty2 00:00:00 gsd-wacom
1947 tty2 00:00:00 gsd-a11y-settin
1948 tty2 00:00:00 gsd-clipboard
1953 tty2 00:00:01 gsd-color
1954 tty2 00:00:00 gsd-datetime
1955 tty2 00:00:00 gsd-housekeepin
1956 tty2 00:00:00 gsd-keyboard
1960 tty2 00:00:00 gsd-media-keys
1965 tty2 00:00:00 gsd-mouse
1985 tty2 00:00:00 gsd-printer
2032 tty2 00:00:00 gsd-disk-utilit
2042 tty2 00:00:02 nautilus-deskto
2085 tty2 00:00:01 ibus-engine-sim
2502 tty2 00:00:06 gnome-software
2508 tty2 00:00:00 update-notifier
4084 tty2 00:00:00 deja-dup-monito
5161 tty2 00:04:11 MainThread
5555 tty2 00:00:13 Privileged Cont
5685 tty2 00:00:08 WebExtensions
5743 tty2 00:05:28 Web Content
5570 tty2 00:04:19 chrome
5577 tty2 00:00:00 cat
5578 tty2 00:00:00 cat
5591 tty2 00:00:00 chrome
5592 tty2 00:00:00 chrome
5602 tty2 00:00:00 nacl_helper
5603 tty2 00:00:00 nacl_helper
5606 tty2 00:00:00 chrome
5634 tty2 00:05:14 chrome
5641 tty2 00:01:24 chrome
5700 tty2 00:00:21 chrome
5717 tty2 00:00:00 chrome
5728 tty2 00:00:00 chrome
```

```
6591 tty2      00:00:00 chrome
6592 tty2      00:00:00 chrome
6602 tty2      00:00:00 nacl_helper
6603 tty2      00:00:00 nacl_helper
6606 tty2      00:00:00 chrome
6634 tty2      00:05:14 chrome
6641 tty2      00:01:24 chrome
6700 tty2      00:00:21 chrome
6717 tty2      00:00:00 chrome
6728 tty2      00:00:00 chrome
7402 tty2      00:00:00 RDD Process
9530 tty2      00:01:14 chrome
9554 tty2      00:00:02 chrome
9922 tty2      00:00:04 chrome
10191 pts/0      00:00:00 sudo
10192 pts/0      00:01:14 wireshark
10200 pts/0      00:00:00 dbus-launch
10379 pts/0      00:00:02 dumpcap
10570 tty2      00:03:26 chrome
11866 tty2      00:00:37 Web Content
12037 pts/0      00:00:00 sudo
12038 pts/0      00:00:00 mn
12081 pts/1      00:00:01 ovs-testcontrol
13461 tty2      00:00:12 chrome
13555 tty2      00:00:24 Web Content
14085 tty2      00:00:01 chrome
14124 tty2      00:00:00 chrome
14139 tty2      00:00:00 Web Content
14975 pts/2      00:00:00 ps
```

This should be the exact same as that seen by the root network namespace:

```
mininet> s1 ps -a
```

```
mininet> s1 ps -a
```

PID	TTY	TIME	CMD
1190	tty1	00:00:00	gnome-session-b
1256	tty1	00:00:24	gnome-shell
1289	tty1	00:00:00	Xwayland
1331	tty1	00:00:00	ibus-daemon
1334	tty1	00:00:00	ibus-dconf
1336	tty1	00:00:00	ibus-x11
1396	tty1	00:00:00	gsd-xsettings
1399	tty1	00:00:00	gsd-a11y-settin
1400	tty1	00:00:00	gsd-clipboard
1401	tty1	00:00:01	gsd-color
1404	tty1	00:00:00	gsd-datetime
1405	tty1	00:00:00	gsd-housekeepin
1406	tty1	00:00:00	gsd-keyboard
1410	tty1	00:00:00	gsd-media-keys
1411	tty1	00:00:00	gsd-mouse
1412	tty1	00:00:00	gsd-power
1413	tty1	00:00:00	gsd-print-notif
1416	tty1	00:00:00	gsd-rfkill
1421	tty1	00:00:00	gsd-screensaver
1424	tty1	00:00:00	gsd-sharing
1428	tty1	00:00:00	gsd-smartcard
1431	tty1	00:00:00	gsd-sound
1433	tty1	00:00:00	gsd-wacom
1447	tty1	00:00:00	ibus-engine-sim
1628	tty2	00:06:55	Xorg
1645	tty2	00:00:00	gnome-session-b
1784	tty2	00:09:10	gnome-shell
1830	tty2	00:00:05	ibus-daemon
1834	tty2	00:00:00	ibus-dconf
1838	tty2	00:00:00	ibus-x11
1911	tty2	00:00:00	gsd-power
1912	tty2	00:00:00	gsd-print-notif
1915	tty2	00:00:00	gsd-rfkill
1919	tty2	00:00:00	gsd-screensaver
1922	tty2	00:00:01	gsd-sharing

1922	tty2	00:00:01	gsd-sharing
1928	tty2	00:00:00	gsd-smartcard
1930	tty2	00:00:00	gsd-sound
1936	tty2	00:00:00	gsd-xsettings
1937	tty2	00:00:00	gsd-wacom
1947	tty2	00:00:00	gsd-a11y-settin
1948	tty2	00:00:00	gsd-clipboard
1953	tty2	00:00:01	gsd-color
1954	tty2	00:00:00	gsd-datetime
1955	tty2	00:00:00	gsd-housekeepin
1956	tty2	00:00:00	gsd-keyboard
1960	tty2	00:00:00	gsd-media-keys
1965	tty2	00:00:00	gsd-mouse
1985	tty2	00:00:00	gsd-printer
2032	tty2	00:00:00	gsd-disk-utilit
2042	tty2	00:00:02	nautilus-deskto
2085	tty2	00:00:01	ibus-engine-sim
2502	tty2	00:00:06	gnome-software
2508	tty2	00:00:00	update-notifier
4084	tty2	00:00:00	deja-dup-monito
5161	tty2	00:04:12	MainThread
5555	tty2	00:00:13	Privileged Cont
5685	tty2	00:00:09	WebExtensions
5743	tty2	00:05:42	Web Content
6570	tty2	00:04:21	chrome
6577	tty2	00:00:00	cat
6578	tty2	00:00:00	cat
6591	tty2	00:00:00	chrome
6592	tty2	00:00:00	chrome
6602	tty2	00:00:00	nacl_helper
6603	tty2	00:00:00	nacl_helper
6606	tty2	00:00:00	chrome
6634	tty2	00:05:14	chrome
6641	tty2	00:01:25	chrome
6700	tty2	00:00:21	chrome
6717	tty2	00:00:00	chrome
6728	tty2	00:00:00	chrome

```

6592 tty2      00:00:00 chrome
6602 tty2      00:00:00 nacl_helper
6603 tty2      00:00:00 nacl_helper
6606 tty2      00:00:00 chrome
6634 tty2      00:05:14 chrome
6641 tty2      00:01:25 chrome
6700 tty2      00:00:21 chrome
6717 tty2      00:00:00 chrome
6728 tty2      00:00:00 chrome
7402 tty2      00:00:00 RDD Process
9530 tty2      00:01:15 chrome
9554 tty2      00:00:02 chrome
9922 tty2      00:00:04 chrome
10191 pts/0     00:00:00 sudo
10192 pts/0     00:01:16 wireshark
10200 pts/0     00:00:00 dbus-launch
10379 pts/0     00:00:02 dumpcap
10570 tty2      00:03:26 chrome
11866 tty2      00:00:38 Web Content
12037 pts/0     00:00:00 sudo
12038 pts/0     00:00:00 mn
12081 pts/1     00:00:01 ovs-testcontrol
13461 tty2      00:00:12 chrome
13555 tty2      00:00:25 Web Content
14085 tty2      00:00:01 chrome
14124 tty2      00:00:00 chrome
14139 tty2      00:00:00 Web Content
15087 pts/4     00:00:00 ps

```

It would be possible to use separate process spaces with Linux containers, but currently Mininet doesn't do that. Having everything run in the "root" process namespace is convenient for debugging, because it allows you to see all of the processes from the console using **ps**, **kill**, etc.

Test connectivity between hosts

Now, verify that you can ping from host 0 to host 1:

mininet> h1 ping -c 1 h2

```

mininet> h1 ping -c 1 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=9.21 ms

--- 10.0.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 9.219/9.219/9.219/0.000 ms

```


If a string appears later in the command with a node name, that node name is replaced by its IP address; this happened for h2.

You should see OpenFlow control traffic. The first host ARPs for the MAC address of the second, which causes a **packet_in** message to go to the controller. The controller then sends a **packet_out** message to flood the broadcast packet to other ports on the switch (in this example, the only other data port). The second host sees the ARP request and sends a reply. This reply goes to the controller, which sends it to the first host and pushes down a flow entry.

Now the first host knows the MAC address of the second, and can send its ping via an ICMP Echo Request. This request, along with its corresponding reply from the second host, both go the controller and result in a flow entry pushed down (along with the actual packets getting sent out).

Repeat the last **ping**:

mininet> h1 ping -c 1 h2

```
mininet> h1 ping -c 1 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=1.32 ms

--- 10.0.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 1.323/1.323/1.323/0.000 ms
```

You should see a much lower **ping** time for the second try (< 100us). A flow entry covering ICMP **ping** traffic was previously installed in the switch, so no control traffic was generated, and the packets immediately pass through the switch.

An easier way to run this test is to use the Mininet CLI built-in **pingall** command, which does an all-pairs **ping**:

mininet> pingall

```
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
```

Run a simple web server and client

Remember that **ping** isn't the only command you can run on a host! Mininet hosts can run any command or application that is available to the underlying Linux system (or VM) and its file system. You can also enter any **bash** command, including job control (**&**, **jobs**, **kill**, etc..)

Next, try starting a simple HTTP server on **h1**, making a request from **h2**, then shutting down the web server:

```
mininet> h1 python -m SimpleHTTPServer 80 &
```

```
mininet> h2 wget -O - h1
```

```
...
```

```
mininet> h1 kill %python
```

```
mininet> h1 python -m SimpleHTTPServer 80 &
mininet> h2 wget -O - h1
wget: invalid option -- '0'
wget: invalid option -- '1'
Usage: wget [OPTION]... [URL]...

Try 'wget --help' for more options.
mininet> h1 kill %python
Serving HTTP on 0.0.0.0 port 80 ...
```

Exit the CLI:

```
mininet> exit
```

```
mininet> exit
*** Stopping 1 controllers
c0
*** Stopping 2 links
..
*** Stopping 1 switches
s1
*** Stopping 2 hosts
h1 h2
*** Done
completed in 10006.398 seconds
```

Cleanup

If Mininet crashes for some reason, clean it up:

\$ sudo mn -c

```
rafatul@rafatul-HP-Notebook:~$ sudo mn -c
[sudo] password for rafatul:
*** Removing excess controllers/ofprotocols/ofdatapaths/pings/noxes
killall controller ofprotocol ofdatapath ping nox_core lt-nox_core ovs-openflowd ovs-controller udpbwtest mnexec ivs 2> /dev/null
killall -9 controller ofprotocol ofdatapath ping nox_core lt-nox_core ovs-openflowd ovs-controller udpbwtest mnexec ivs 2> /dev/null
pkill -9 -f "sudo mnexec"
*** Removing junk from /tmp
rm -f /tmp/vconn* /tmp/vlogs* /tmp/*.out /tmp/*.log
*** Removing old X11 tunnels
*** Removing excess kernel datapaths
ps ax | egrep -o 'dp[0-9]+' | sed 's/dp/nl:/'
*** Removing OVS datapaths
ovs-vsctl --timeout=1 list-br
ovs-vsctl --timeout=1 list-br
*** Removing all links of the pattern foo-ethX
ip link show | egrep -o '([-.[:alnum:]]+-eth[[:digit:]]+)'
ip link show
*** Killing stale mininet node processes
pkill -9 -f mininet:
*** Shutting down stale tunnels
pkill -9 -f Tunnel=Ethernet
pkill -9 -f .ssh/mn
rm -f ~/.ssh/mn/*
*** Cleanup complete.
rafatul@rafatul-HP-Notebook:~$
```